# HoneyHive POC Plan

*Two-stage proof of concept for agent observability and evaluation*

**Author:** Amine Bellajdel
**Date:** April 30, 2026

## Framing

The platform team is excited to try HoneyHive. The challenge is that two goals are in tension. A POC that minimizes procurement and legal involvement is necessarily small — narrow scope, short timeline, low dollar amount. A POC that positions HoneyHive for a strong ACV (Annual Contract Value) land needs to be broad enough to demonstrate value across multiple agents and to involve Responsible AI directly. Push too hard on procurement-light and the deal anchors at a number too small to matter. Push too hard on ACV-strong and the deal stalls in procurement for six to nine months.

The resolution is a two-stage POC. Stage 1 is procurement-light by design and produces technical validation. Stage 2 is paid, narrowly scoped, and produces the RAI evidence package that justifies a real ACV land in stage 3. The first stage's job is not to make money. Its job is to manufacture the evidence and the internal champion that make the second stage an easy yes.

## The customer's actual situation

A clarification that shapes the entire plan: the bank already has 5 agents running in production. They reached production before Responsible AI tightened controls. New use cases are queued every month, and none of them can ship until platform demonstrates satisfactory quality controls.

This means HoneyHive's value proposition is not only about unblocking future agents. It is also, urgently, about giving RAI visibility into the 5 agents already in production — which today run without trace-level instrumentation. Those 5 agents are a current liability for RAI, not a future one. The POC should address both: remediate the existing 5, then enable the queued 45.

## Stage 1 — Technical validation (30 days, procurement-light)

### Goal

Prove that HoneyHive's trace data and tool-misuse evaluator work against real production traffic from one of the 5 in-production agents. Build technical confidence inside the platform team before any RAI conversation.

### Scope

One agent of the bank's choice — recommended: the fraud triage agent, given its self-contained tool surface and clear policy boundaries. Self-hosted deployment in the bank's AWS VPC (Virtual Private Cloud) from day one — no customer data leaves the bank's environment.

Shadow-mode instrumentation only. Traces are captured and evaluated but no production behavior changes. This minimizes the security review surface and gets to first traces in days, not weeks.

### Audience

Platform team only. RAI is informed but not engaged as a reviewer yet. The platform team's confidence is the prerequisite for a credible RAI conversation in stage 2.

**Commercial structure**

Paid POC under $50,000, scoped to one agent for 30 days. This fits comfortably under the AI Platform Director's discretionary spend authority and triggers only the light-touch procurement review (~2 weeks, short-form contract). Self-hosted deployment avoids the AI-specific addendum that adds 4–6 weeks to standard AI vendor review.

Why paid rather than free: paid pilots get attention. Free trials get neglected, and a stalled free trial does not produce the champion needed for stage 2.

**Success criteria**

Defined jointly with the platform team in the kickoff meeting, before any code is written. Recommended criteria:

1. Traces from the chosen agent are visible in HoneyHive within 5 business days of contract signature.

2. The tool-misuse evaluator runs against production traffic and flags at least one historical trace pattern the platform team recognizes as concerning.

3. The on-call engineer can debug a production incident using HoneyHive traces faster than with current tooling. Measured against one real incident or one synthetic exercise.

4. Cost, latency, and quality dashboards are populated and useful to the platform team.

**Exit criteria**

Stage 1 is complete when all four success criteria are met. If any criterion is not met, both sides discuss whether to extend, modify, or end. No automatic conversion to stage 2 — this is a real decision point.

## Stage 2 — RAI evidence pilot (60 days, paid)

**Goal**

Produce the first RAI-approved evidence package for an in-production agent. Establish the template and review cadence that will then apply to the other 4 in-production agents and the 45 queued.

**Scope**

Same agent as stage 1, plus one additional in-production agent of RAI's choice. Two agents is the minimum to demonstrate that the instrumentation pattern and evaluator framework are reusable across domains and frameworks. Both agents covered with the full evaluator suite: hallucinations, PII (Personally Identifiable Information) leakage, tool misuse.

**Audience**

Platform team plus RAI as a defined reviewer. RAI participates in the kickoff to define evaluator thresholds, reviews evidence at day 30, and signs off (or asks for changes) at day 60. The day-60 review is the gating event for stage 3.

**Commercial structure**

Paid pilot, sized to remain under $50,000 if possible to retain light-touch procurement, or up to ~$75,000 if scope demands it. Above $50,000 triggers standard procurement (~3 months, expedited path possible) but the platform team will have already-signed stage 1 evidence to accelerate review.

**Success criteria**

Defined jointly with the platform team and RAI in the kickoff meeting. Recommended criteria:

1. Both agents fully instrumented with all three evaluator types (hallucinations, PII, tool misuse) running continuously against production traffic.

2. RAI receives evidence in a format that maps to the bank's existing MRM (Model Risk Management) framework under SR 11-7 (Federal Reserve Supervisory Letter 11-7).

3. RAI signs off that the evidence package is sufficient to support ongoing approval of these two agents, and that the same template, applied to other agents, would meet RAI's review standard.

4. The platform team has a documented runbook for instrumenting additional agents, with time-to-instrumentation under 5 business days per agent.

### Exit criteria

Stage 2 is complete at day 60 with RAI's go/no-go decision. A go-decision is the trigger for stage 3 negotiations. A no-go decision triggers a joint review of what would need to change before broader rollout.

## Stage 3 — The ACV land

### Goal

Convert the validated POC into a multi-year platform commitment that covers both the remediation need (the 5 agents already in production) and the enablement need (the 45 queued and any future use cases).

### Scope

Two work streams in one contract:

**Remediation.** Instrument the remaining 3 in-production agents within 90 days of contract signature. Bring all 5 under the same evaluator framework and review cadence agreed with RAI in stage 2. This addresses the open audit risk RAI is most concerned about today.

**Enablement.** Standardize HoneyHive as the agent observability and evaluation layer for all future use cases. Establish per-agent onboarding playbooks, CI/CD (Continuous Integration / Continuous Deployment) integration patterns, and self-service workflows for the use-case teams to instrument their own agents at the same standard.

### Commercial structure

Multi-year contract priced per agent covered, with volume tiers as the customer scales from 5 to 50. Per-agent pricing maps directly to the customer's internal narrative and produces a predictable line item that finance can budget. Includes self-hosted deployment, support tier, and access to all three product surfaces (observability, evaluation, prompt management).

Triggers full enterprise procurement (~6–9 months) but the path is materially de-risked by stage 1 and stage 2: existing security review, RAI sign-off already in hand, deployment posture proven, champion already created.

## Risks and mitigations

| Risk | Mitigation |
|---|---|
| Stage 1 stalls in security review even at low dollar amount. | Self-hosted deployment from day one removes the data-handling questions that typically slow AI vendors. CISO security questionnaire begins on day 1 of stage 1. |
| RAI declines to engage at stage 2 kickoff. | Stage 1 produces concrete artifacts (traces, evaluator hits) that the platform team brings to RAI directly. The ask is narrow: review one agent's evidence, not approve a vendor. |
| Success criteria drift mid-POC. | Criteria are written down in the kickoff and signed off by all parties before any work begins. Mid-POC changes require joint sign-off. |
| Procurement dollar threshold creep. | Stage 1 is hard-capped at $50,000 to stay under the Director's discretionary authority. Stage 2 is sized deliberately to avoid the AI-specific addendum where possible. |
| POC becomes permanent — never converts to stage 3. | Each stage has an explicit exit decision point, not automatic continuation. Stage 3 negotiation begins on day 45 of stage 2, in parallel with the final review. |

## Summary

| | Stage 1 | Stage 2 | Stage 3 |
|---|---|---|---|
| **Duration** | 30 days | 60 days | Multi-year |
| **Scope** | 1 agent, shadow mode | 2 agents, full evaluators | All 5 in-production + framework for future agents |
| **Audience** | Platform team only | Platform + RAI as reviewer | Platform + RAI + business line owners |
| **Commercial** | Paid, under $50K | Paid, $50K–$75K | Multi-year, per-agent pricing with volume tiers |
| **Procurement path** | Light-touch (~2 weeks, director discretion) | Standard, expedited (~3 months) — possibly light-touch if under $50K | Full enterprise (~6–9 months) |
| **Gating event** | Platform team confidence | RAI sign-off on evidence | Contract signature |

## How this plan resolves the tension

Stage 1 stays small enough to slip past procurement quickly and produces concrete technical artifacts. Stage 2 builds on those artifacts to engage RAI in a narrow, defensible way that produces the evidence the bank actually needs to unblock its agent program. Stage 3 is the real ACV deal, and by the time it begins, every major objection — security, RAI, deployment posture, internal champion — has already been answered.

The deal does not anchor at the small POC number, because stage 1 is explicitly framed as technical validation, not value validation. And the deal does not stall in procurement, because stage 1 is deliberately sized to avoid the long-form review path. The two goals stop being in tension once the POC is sequenced rather than collapsed into a single event.